



Sardar Patel Institute of Technology, Mumbai

Department of Computer Science Engineering

B.E. Sem-VII- PE-IV (2024-2025)

IT 24 - AI in Healthcare

Experiment5: Data Transformation and preparation for analysis of healthcare data

Name: Adwait Purao

Date: 7/10/24

Objective:

To understand and apply data transformation techniques for preparing healthcare data for analysis.

Link to Dataset :

<https://www.kaggle.com/datasets/prasad22/healthcare-dataset/data>

Link to Notebook:

<https://colab.research.google.com/drive/1w8aKUO0BfcGsN6VhiqxGmCIDwM6mJIG3?usp=sharing>

1.Introduction to Healthcare Data:

Dataset Information:

Each column provides specific information about the patient, their admission, and the healthcare services provided, making this dataset suitable for various data analysis and

modeling tasks in the healthcare domain. Here's a brief explanation of each column in the dataset -

- **Name:** This column represents the name of the patient associated with the healthcare record.
- **Age:** The age of the patient at the time of admission, expressed in years.
- **Gender:** Indicates the gender of the patient, either "Male" or "Female."
- **Blood Type:** The patient's blood type, which can be one of the common blood types (e.g., "A+", "O-", etc.).
- **Medical Condition:** This column specifies the primary medical condition or diagnosis associated with the patient, such as "Diabetes," "Hypertension," "Asthma," and more.
- **Date of Admission:** The date on which the patient was admitted to the healthcare facility.
- **Doctor:** The name of the doctor responsible for the patient's care during their admission.
- **Hospital:** Identifies the healthcare facility or hospital where the patient was admitted.
- **Insurance Provider:** This column indicates the patient's insurance provider, which can be one of several options, including "Aetna," "Blue Cross," "Cigna," "UnitedHealthcare," and "Medicare."
- **Billing Amount:** The amount of money billed for the patient's healthcare services during their admission. This is expressed as a floating-point number.
- **Room Number:** The room number where the patient was accommodated during their admission.
- **Admission Type:** Specifies the type of admission, which can be "Emergency," "Elective," or "Urgent," reflecting the circumstances of the admission.
- **Discharge Date:** The date on which the patient was discharged from the healthcare facility, based on the admission date and a random number of days within a realistic range.
- **Medication:** Identifies a medication prescribed or administered to the patient during their admission. Examples include "Aspirin," "Ibuprofen," "Penicillin," "Paracetamol," and "Lipitor."
- **Test Results:** Describes the results of a medical test conducted during the patient's admission. Possible values include "Normal," "Abnormal," or "Inconclusive," indicating the outcome of the test.

2.Understanding the Dataset

Dataset Description:

```
import pandas as pd

def display_step_header(step_name):

    print(f"\n{'='*20} {step_name} {'='*20}")
```

```
def load_data(file_path):  
  
    display_step_header("1. Data Loading")  
  
    df = pd.read_csv(file_path)  
  
    print("Dataset Overview:")  
  
    print(df.info())  
  
    print("\nSample of the dataset:")  
  
    print(df.head())  
  
  
    print("\nBasic statistics of numerical columns:")  
  
    print(df.describe())  
  
  
    # Display value counts for categorical columns  
    categorical_columns = df.select_dtypes(include=['object']).columns  
    for col in categorical_columns:  
        print(f"\nValue counts for {col}:")  
        print(df[col].value_counts().head())  
  
    return df  
  
df = load_data('./healthcare_dataset.csv')
```

===== 1. Data Loading =====

Dataset Overview:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 55500 entries, 0 to 55499

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	Name	55500 non-null	object
1	Age	55500 non-null	int64
2	Gender	55500 non-null	object
3	Blood Type	55500 non-null	object
4	Medical Condition	55500 non-null	object
5	Date of Admission	55500 non-null	object
6	Doctor	55500 non-null	object
7	Hospital	55500 non-null	object
8	Insurance Provider	55500 non-null	object
9	Billing Amount	55500 non-null	float64
10	Room Number	55500 non-null	int64
11	Admission Type	55500 non-null	object
12	Discharge Date	55500 non-null	object
13	Medication	55500 non-null	object
14	Test Results	55500 non-null	object

dtypes: float64(1), int64(2), object(12)

memory usage: 6.4+ MB

None

Sample of the dataset:

	Name	Age	Gender	Blood Type	Medical Condition	Date of Admission	\
0	Bobby JacksOn	30	Male	B-	Cancer	2024-01-31	
1	Leslie TErRy	62	Male	A+	Obesity	2019-08-20	
2	DaNnY sMitH	76	Female	A-	Obesity	2022-09-22	
3	andrEw waTtS	28	Female	O+	Diabetes	2020-11-18	
4	adrIENNE bEll	43	Female	AB+	Cancer	2022-09-19	

	Doctor	Hospital	Insurance Provider	\
0	Matthew Smith	Sons and Miller	Blue Cross	
1	Samantha Davies	Kim Inc	Medicare	
2	Tiffany Mitchell	Cook PLC	Aetna	
3	Kevin Wells	Hernandez Rogers and Vang,	Medicare	
4	Kathleen Hanna	White-White	Aetna	

	Billing Amount	Room Number	Admission Type	Discharge Date	Medication	\
0	18856.281306	328	Urgent	2024-02-02	Paracetamol	
1	33643.327287	265	Emergency	2019-08-26	Ibuprofen	
2	27955.096079	205	Emergency	2022-10-07	Aspirin	
3	37909.782410	450	Elective	2020-12-18	Ibuprofen	
4	14238.317814	458	Urgent	2022-10-09	Penicillin	

Test Results

0	Normal
1	Inconclusive
2	Normal
3	Abnormal
4	Abnormal

Basic statistics of numerical columns:

	Age	Billing Amount	Room Number
count	55500.000000	55500.000000	55500.000000

Basic statistics of numerical columns:

	Age	Billing Amount	Room Number
count	55500.000000	55500.000000	55500.000000
mean	51.539459	25539.316097	301.134829
std	19.602454	14211.454431	115.243069
min	13.000000	-2008.492140	101.000000
25%	35.000000	13241.224652	202.000000
50%	52.000000	25538.069376	302.000000
75%	68.000000	37820.508436	401.000000
max	89.000000	52764.276736	500.000000

Value counts for Name:

Name

DAvId muNoZ	3
SOnYa aDams	2
terRY gONZaLeZ	2
JaCKsON BARbeR	2
doNALD aViLA	2

Name: count, dtype: int64

Value counts for Gender:

Gender

Male	27774
Female	27726

Name: count, dtype: int64

Value counts for Blood Type:

Blood Type

A-	6969
A+	6956
AB+	6947
AB-	6945
B+	6945

Name: count, dtype: int64

Value counts for Medical Condition:

Medical Condition

Arthritis	9308
Diabetes	9304
Hypertension	9245
Obesity	9231
Cancer	9227

Name: count, dtype: int64

Value counts for Date of Admission:

Date of Admission

2024-03-16	50
2022-07-24	49
2020-10-22	49
2021-12-28	48
2021-01-03	48

Name: count, dtype: int64

Value counts for Doctor:

Doctor

Michael Smith	27
Robert Smith	22
John Smith	22
Michael Johnson	20

Value counts for Doctor:

Doctor

Michael Smith 27

Robert Smith 22

John Smith 22

Michael Johnson 20

James Smith 20

Name: count, dtype: int64

Value counts for Hospital:

Hospital

LLC Smith 44

Ltd Smith 39

Johnson PLC 38

Smith Ltd 37

Smith PLC 36

Name: count, dtype: int64

Value counts for Insurance Provider:

Insurance Provider

Cigna 11249

Medicare 11154

UnitedHealthcare 11125

Blue Cross 11059

Aetna 10913

Name: count, dtype: int64

Value counts for Admission Type:

Admission Type

Elective 18655

Urgent 18576

Emergency 18269

Name: count, dtype: int64

Value counts for Discharge Date:

Discharge Date

2020-03-15 53

2021-12-13 51

2020-12-02 51

2023-04-29 51

2020-08-11 50

Name: count, dtype: int64

Value counts for Medication:

Medication

Lipitor 11140

Ibuprofen 11127

Aspirin 11094

Paracetamol 11071

Penicillin 11068

Name: count, dtype: int64

Value counts for Test Results:

Test Results

Abnormal 18627

Normal 18517

Inconclusive 18356

Name: count, dtype: int64

3. Handling Missing data in the dataset:

Methods of Handling Missing data

```
def handle_missing_data(df):  
  
    display_step_header("2. Missing Data Handling")  
  
    print("Missing values before handling:")  
  
    print(df.isnull().sum())  
  
    # For numerical columns, fill with median  
  
    numeric_columns = df.select_dtypes(include=['int64',  
'float64']).columns  
  
    for col in numeric_columns:  
        df[col].fillna(df[col].median(), inplace=True)  
  
    # For categorical columns, fill with mode  
  
    categorical_columns = df.select_dtypes(include=['object']).columns  
  
    for col in categorical_columns:  
        df[col].fillna(df[col].mode()[0], inplace=True)  
  
    print("\nMissing values after handling:")  
  
    print(df.isnull().sum())  
  
    return df  
  
df = handle_missing_data(df)
```

```

===== 2. Missing Data Handling =====
Missing values before handling:
Name          0
Age           0
Gender        0
Blood Type    0
Medical Condition 0
Date of Admission 0
Doctor        0
Hospital      0
Insurance Provider 0
Billing Amount 0
Room Number   0
Admission Type 0
Discharge Date 0
Medication    0
Test Results  0
dtype: int64

Cipython-Input-12-91a262dc2575>:9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

df[col].fillna(df[col].median(), inplace=True)
Cipython-Input-12-91a262dc2575>:14: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

df[col].fillna(df[col].mode()[0], inplace=True)

Missing values after handling:
Name          0
Age           0
Gender        0
Blood Type    0
Medical Condition 0
Date of Admission 0
Doctor        0
Hospital      0
Insurance Provider 0
Billing Amount 0
Room Number   0
Admission Type 0
Discharge Date 0
Medication    0
Test Results  0
dtype: int64

```

4.Categorical Data Encoding

```

from sklearn.preprocessing import LabelEncoder

def encode_categorical_data(df):

    display_step_header("3. Categorical Data Encoding")

    label_encoders = {}

    categorical_columns = df.select_dtypes(include=['object']).columns

    # Create a sample dataframe to show transformations

    sample_df = pd.DataFrame()

    for col in categorical_columns:

        label_encoders[col] = LabelEncoder()

        df[f'{col}_encoded'] = label_encoders[col].fit_transform(df[col])

```



```

        # Get unique categories and their encodings

        unique_categories = df[col].unique()[:3] # Take only first 3 unique
values
        unique_encodings = [label_encoders[col].transform([cat])[0] for cat
in unique_categories]

        # Ensure sample_df has the correct length

        sample_df = pd.DataFrame({col: unique_categories}) # Create a
DataFrame with the correct length

        sample_df[f'{col}_encoded'] = unique_encodings

        print(f"\nEncoding example for {col} (showing first 3 categories):")

        for orig, enc in zip(unique_categories, unique_encodings):

            print(f"{orig} -> {enc}")

        print("\nSample of original and encoded data:")

        print(sample_df.to_string(index=False))

    return df, label_encoders

df, label_encoders = encode_categorical_data(df)

```

===== 3. Categorical Data Encoding =====

Encoding example for Name (showing first 3 categories):

Bobby JacksOn -> 3068

Leslie TErRy -> 15211

DaNnY sMitH -> 6476

Encoding example for Gender (showing first 3 categories):

Male -> 1

Female -> 0

Encoding example for Blood Type (showing first 3 categories):

B- -> 5

A+ -> 0

A- -> 1

Encoding example for Medical Condition (showing first 3 categories):

Cancer -> 2

Obesity -> 5

Diabetes -> 3

Encoding example for Date of Admission (showing first 3 categories):

2024-01-31 -> 1729

2019-08-20 -> 104

2022-09-22 -> 1233

Encoding example for Doctor (showing first 3 categories):

Matthew Smith -> 26612

Samantha Davies -> 33648

Tiffany Mitchell -> 37828

Encoding example for Hospital (showing first 3 categories):

Sons and Miller -> 29933

Kim Inc -> 16012

Cook PLC -> 5473

Encoding example for Insurance Provider (showing first 3 categories):

Blue Cross -> 1

Medicare -> 3

Aetna -> 0

Encoding example for Admission Type (showing first 3 categories):

Urgent -> 2

Emergency -> 1

Elective -> 0

Encoding example for Discharge Date (showing first 3 categories):

2024-02-02 -> 1730

2019-08-26 -> 109

2022-10-07 -> 1247

Encoding example for Medication (showing first 3 categories):

Paracetamol -> 3

Ibuprofen -> 1

Aspirin -> 0

Encoding example for Test Results (showing first 3 categories):

Normal -> 2

Inconclusive -> 1

```
Encoding example for Test Results (showing first 3 categories):  
Normal -> 2  
Inconclusive -> 1  
Abnormal -> 0
```

```
Sample of original and encoded data:  
Test Results  Test Results_encoded  
    Normal                2  
Inconclusive                1  
    Abnormal                0
```

5. Outlier Detection and Treatment

```
def handle_outliers(df, numeric_columns):  
  
    display_step_header("4. Outlier Detection and Treatment")  
  
    for col in numeric_columns:  
  
        Q1 = df[col].quantile(0.25)  
  
        Q3 = df[col].quantile(0.75)  
  
        IQR = Q3 - Q1  
  
        lower_bound = Q1 - 1.5 * IQR  
  
        upper_bound = Q3 + 1.5 * IQR  
  
        # Print outlier information  
  
        outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]  
  
        print(f"\nOutliers in {col}:")  
  
        print(f"Number of outliers: {len(outliers)}")  
  
        print(f"Percentage                of                outliers:  
{ (len(outliers)/len(df)) * 100: .2f}%")  
  
        print(f"Lower bound: {lower_bound: .2f}")  
  
        print(f"Upper bound: {upper_bound: .2f}")  
  
        # Cap the outliers
```

```
df[f'{col}_cleaned'] = df[col].clip(lower=lower_bound,
upper=upper_bound)

# Display statistics before and after
print(f"\nStatistics for {col} before and after outlier treatment:")
print(pd.DataFrame({
    'Original': df[col].describe(),
    'Cleaned': df[f'{col}_cleaned'].describe()
}))

return df

numeric_columns = ['Age', 'Billing Amount']
df = handle_outliers(df, numeric_columns)
```

===== 4. Outlier Detection and Treatment =====

Outliers in Age:

Number of outliers: 0

Percentage of outliers: 0.00%

Lower bound: -14.50

Upper bound: 117.50

Statistics for Age before and after outlier treatment:

	Original	Cleaned
count	55500.000000	55500.000000
mean	51.539459	51.539459
std	19.602454	19.602454
min	13.000000	13.000000
25%	35.000000	35.000000
50%	52.000000	52.000000
75%	68.000000	68.000000
max	89.000000	89.000000

Outliers in Billing Amount:

Number of outliers: 0

Percentage of outliers: 0.00%

Lower bound: -23627.70

Upper bound: 74689.43

Statistics for Billing Amount before and after outlier treatment:

	Original	Cleaned
count	55500.000000	55500.000000
mean	25539.316097	25539.316097
std	14211.454431	14211.454431
min	-2008.492140	-2008.492140
25%	13241.224652	13241.224652
50%	25538.069376	25538.069376
75%	37820.508436	37820.508436
max	52764.276736	52764.276736

6. Feature Scaling and Normalization

```
from sklearn.preprocessing import StandardScaler
```

```
def scale_features(df, numeric_columns):
```

```
    display_step_header("5. Feature Scaling")
```

```
    scaler = StandardScaler()
```

```
scaled_columns = [col + '_scaled' for col in numeric_columns]

df[scaled_columns] = scaler.fit_transform(df[numeric_columns])


for original, scaled in zip(numeric_columns, scaled_columns):

    print(f"\nScaling results for {original}:")

    print(pd.DataFrame({

        'Original': df[original].describe(),

        'Scaled': df[scaled].describe()

    })))


return df, scaler


df, scaler = scale_features(df, numeric_columns)
```

===== 5. Feature Scaling =====

Scaling results for Age:

	Original	Scaled
count	55500.000000	5.550000e+04
mean	51.539459	7.732753e-17
std	19.602454	1.000009e+00
min	13.000000	-1.966071e+00
25%	35.000000	-8.437519e-01
50%	52.000000	2.349424e-02
75%	68.000000	8.397259e-01
max	89.000000	1.911030e+00

Scaling results for Billing Amount:

	Original	Scaled
count	55500.000000	5.550000e+04
mean	25539.316097	5.703546e-17
std	14211.454431	1.000009e+00
min	-2008.492140	-1.938440e+00
25%	13241.224652	-8.653725e-01
50%	25538.069376	-8.772730e-05
75%	37820.508436	8.641834e-01
max	52764.276736	1.915723e+00

7. Feature Engineering

- Create new features (e.g., BMI from weight and height).

```
def engineer_features(df):  
  
    display_step_header("6. Feature Engineering")  
  
    # Age groups  
  
    df['Age_Group'] = pd.cut(df['Age'], bins=[0, 18, 35, 50, 65,  
100],  
  
                             labels=['0-18', '19-35', '36-50', '51-  
65', '65+'])  
  
    print("\nAge Group Distribution:")  
  
    print(df['Age_Group'].value_counts())
```

```
# Chronic condition flag

df['Is_Chronic'] = df['Medical Condition'].apply(lambda x: 1 if
'chronic' in str(x).lower() else 0)

print("\nChronic Condition Distribution:")

print(df['Is_Chronic'].value_counts(normalize=True))


# Cost per age

df['Cost_Per_Age'] = df['Billing Amount'] / df['Age']

print("\nCost Per Age Statistics:")

print(df['Cost_Per_Age'].describe())


return df


df = engineer_features(df)
```



```

===== 6. Feature Engineering =====

Age Group Distribution:
Age_Group
65+      16250
19-35    13644
51-65    12417
36-50    12301
0-18      888
Name: count, dtype: int64

Chronic Condition Distribution:
Is_Chronic
0      1.0
Name: proportion, dtype: float64

Cost Per Age Statistics:
count      55500.000000
mean        596.195820
std         467.843388
min         -49.140205
25%         257.961738
50%         495.787257
75%         787.818365
max         3886.670220
Name: Cost_Per_Age, dtype: float64

```

8. Data Preparation for Machine Learning Models:

- Train-Test Split

```

from sklearn.model_selection import train_test_split

def prepare_for_ml(df, target_column='Billing Amount'):

    display_step_header("7. Preparing Data for Machine Learning")

    numeric_features = df.select_dtypes(include=['int64',
'float64']).columns

    categorical_features = [col for col in df.columns if '_encoded'
in col]

```

```
features = list(numeric_features) + categorical_features

features = [f for f in features if f != target_column]


X = df[features]

y = df[target_column]


X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)


print("Selected features:")

for f in features:

    print(f"- {f}")


print(f"\nTarget variable statistics:")

print(y.describe())


print("\nData split sizes:")

print(f"Training set: {X_train.shape}")

print(f"Test set: {X_test.shape}")


return X_train, X_test, y_train, y_test


X_train, X_test, y_train, y_test = prepare_for_ml(df)
```

===== 7. Preparing Data for Machine Learning =====

Selected features:

- Age
- Room Number
- Name_encoded
- Gender_encoded
- Blood Type_encoded
- Medical Condition_encoded
- Date of Admission_encoded
- Doctor_encoded
- Hospital_encoded
- Insurance Provider_encoded
- Admission Type_encoded
- Discharge Date_encoded
- Medication_encoded
- Test Results_encoded
- Age_cleaned
- Billing Amount_cleaned
- Age_scaled
- Billing Amount_scaled
- Is_Chronic
- Cost_Per_Age
- Name_encoded
- Gender_encoded
- Blood Type_encoded
- Medical Condition_encoded
- Date of Admission_encoded
- Doctor_encoded
- Hospital_encoded
- Insurance Provider_encoded
- Admission Type_encoded
- Discharge Date_encoded
- Medication_encoded
- Test Results_encoded

Target variable statistics:

count	55500.000000
mean	25539.316097
std	14211.454431
min	-2008.492140
25%	13241.224652
50%	25538.069376
75%	37820.508436
max	52764.276736

Name: Billing Amount, dtype: float64

Data split sizes:

Training set: (44400, 32)

Test set: (11100, 32)

Conclusion:

In this assignment, I learned how to effectively handle real-world data by applying essential preprocessing steps such as handling missing values, encoding categorical variables, detecting and treating outliers, scaling features, and engineering new features. I also gained insights into how these steps impact the data's quality and readiness for machine learning. Through each step, I enhanced my understanding of preparing datasets for analysis and model building, ensuring the data is clean, consistent, and well-structured for successful predictions.